


a TabM: parameter-sharing ensemble


Obtain several neural predictors while sharing most of their weights.



32 members share parameters; they are not 32 independent fits.

b TabPFN-3.5: prediction from context

Condition a pretrained transformer on eligible labelled training rows.




ARROW MAE 0.199; size-held-out MAE 1.318.

Native intervals under-cover; input recipes differ across partitions.

c Frozen MoLFormer + ExtraTrees

Change the representation while retaining the original endpoint family.



ARROW MAE 0.234 with responses; 0.348 without.

This control retains the bounded tree endpoint.