


a TabM: a parameter-sharing ensemble

Use member-specific factors within layers that share their main weights.




External MAE 1.132 · Size-held-out MAE 1.389

32 members share parameters; these are not 32 independent training runs.

b TabPFN-3.5: prediction from context

Condition a pretrained transformer on eligible labelled training rows.




ARROW MAE 0.199 · Size-held-out MAE 1.318

Native intervals under-cover; input recipes differ across partitions.

c Frozen MoLFormer + ExtraTrees

Change the representation while keeping the tree endpoint family.



ARROW MAE 0.234 · Without responses: 0.348

The endpoint remains bounded by its stored leaf values.